

Deep Reinforcement Learning

Deep Q-Networks

Carlos Morato

Used Materials

- Disclaimer:** Much of the material and slides for this lecture were borrowed from Rich Sutton's RL class and David Silver's Deep RL tutorial

Components of an RL Agent

- An RL agent may include one or more of these components:
 - **Policy**: agent's behavior function
 - **Value function**: how good is each state and/or action
 - **Model**: agent's representation of the environment
- A policy is the agent's behavior
- It is a map from state to action:
 - **Deterministic** policy: $a = \pi(s)$
 - **Stochastic** policy: $\pi(a|s) = P[a|s]$

Review: Value Function

- A value function is a prediction of **future reward**
 - How much reward will I get from action a in state s ?
- Q-value function gives **expected total reward**
 - from state s and action a
 - under policy π
 - with discount factor γ

$$Q^\pi(s, a) = \mathbb{E} [r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots \mid s, a]$$

- Value functions decompose into a **Bellman equation**

$$Q^\pi(s, a) = \mathbb{E}_{s', a'} [r + \gamma Q^\pi(s', a') \mid s, a]$$

Optimal Value Function

- An optimal value function is the **maximum achievable value**

$$Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a) = Q^{\pi^*}(s, a)$$

- Once we have Q^* , the agent can act optimally

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$

- Formally, optimal values decompose into a **Bellman equation**

$$Q^*(s, a) = \mathbb{E}_{s'} \left[r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]$$

Optimal Value Function

- An optimal value function is the **maximum achievable value**

$$Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a) = Q^{\pi^*}(s, a)$$

- Formally, optimal values decompose into a **Bellman equation**

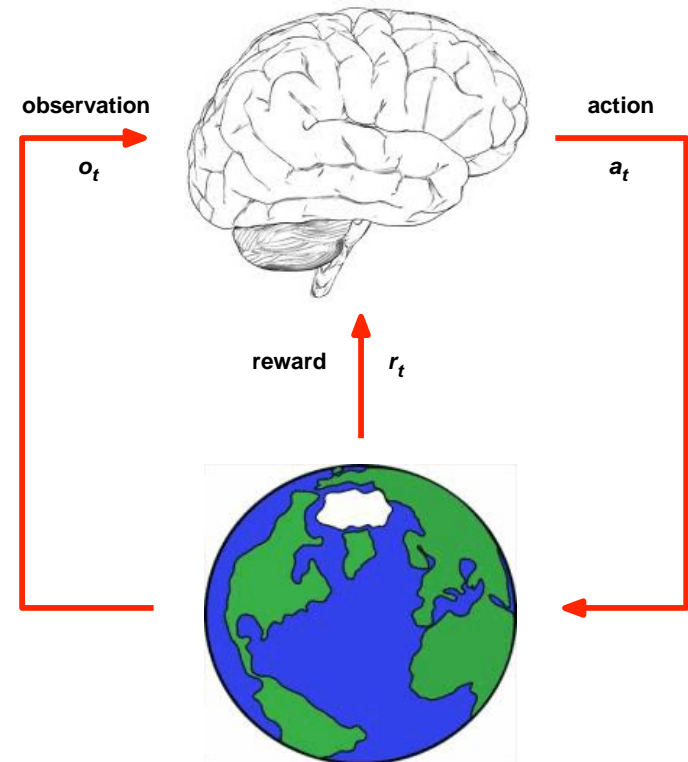
$$Q^*(s, a) = \mathbb{E}_{s'} \left[r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]$$

- **Informally**, optimal value maximizes over all decisions

$$\begin{aligned} Q^*(s, a) &= r_{t+1} + \gamma \max_{a_{t+1}} r_{t+2} + \gamma^2 \max_{a_{t+2}} r_{t+3} + \dots \\ &= r_{t+1} + \gamma \max_{a_{t+1}} Q^*(s_{t+1}, a_{t+1}) \end{aligned}$$

Model

- Model is learned from **experience**
- Acts as proxy for environment
- Planner interacts with model, e.g. using look-ahead search



Approaches to RL

- **Value-based RL** (this is what we have looked at so far)
 - Estimate the optimal value function $Q^*(s,a)$
 - This is the maximum value achievable under any policy
- **Policy-based RL**
 - Search directly for the optimal policy π^*
 - This is the policy achieving maximum future reward
- **Model-based RL**
 - Build a model of the environment
 - Plan (e.g. by look-ahead) using model
- Let us revisit value-based RL.

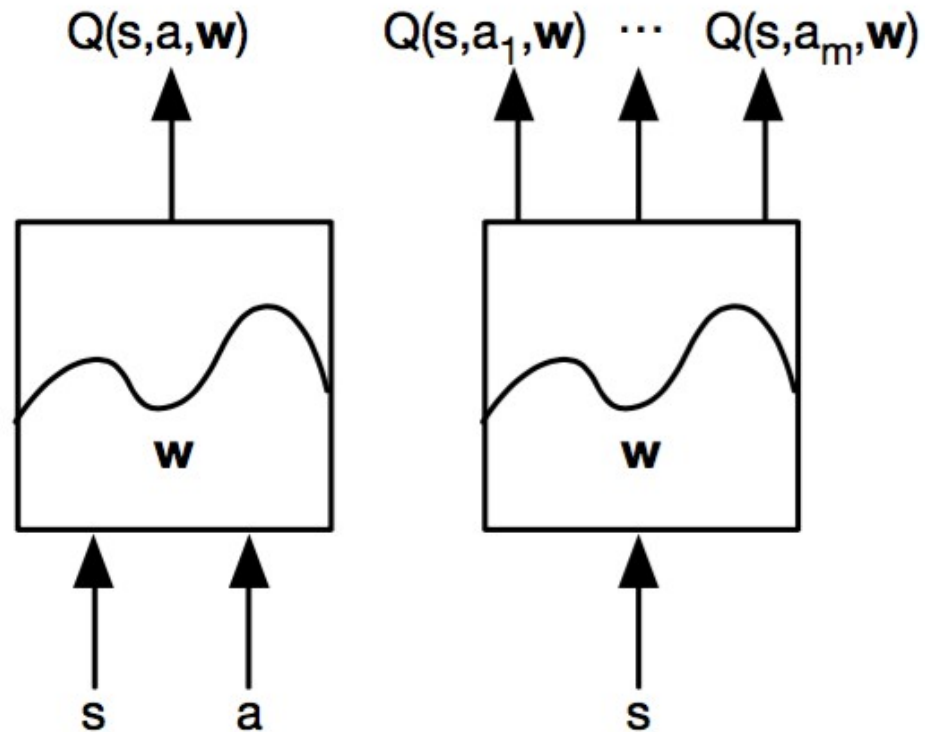
Deep Reinforcement Learning

- Use deep neural networks to represent
 - Value function
 - Policy
 - Model
- Optimize loss function by stochastic gradient descent (SGD)

Deep Q-Networks (DQNs)

- Represent value function by Q-network with weights w

$$Q(s, a, \mathbf{w}) \approx Q^*(s, a)$$



Q-Learning

- Optimal Q-values should obey Bellman equation

$$Q^*(s, a) = \mathbb{E}_{s'} \left[r + \gamma \max_{a'} Q(s', a')^* \mid s, a \right]$$

- Treat right-hand $r + \gamma \max_{a'} Q(s', a', \mathbf{w})$ as a target
- Minimize MSE loss by stochastic gradient descent

$$l = \left(r + \gamma \max_a Q(s', a', \mathbf{w}) - Q(s, a, \mathbf{w}) \right)^2$$

- Remember VFA lecture: Minimize mean-squared error between the true action-value function $q_\pi(S, A)$ and the approximate Q function:

$$J(\mathbf{w}) = \mathbb{E}_\pi \left[(q_\pi(S, A) - \hat{q}(S, A, \mathbf{w}))^2 \right]$$

Q-Learning

- Minimize MSE loss by stochastic gradient descent

$$l = \left(r + \gamma \max_a Q(s', a', \mathbf{w}) - Q(s, a, \mathbf{w}) \right)^2$$

- Converges to Q^* using **table lookup representation**
- But diverges using neural networks due to:
 - Correlations between samples
 - Non-stationary targets

DQNs: Experience Replay

- To remove correlations, build data-set from agent's own experience

s_1, a_1, r_2, s_2	\rightarrow s, a, r, s'
s_2, a_2, r_3, s_3	
s_3, a_3, r_4, s_4	
...	
$s_t, a_t, r_{t+1}, s_{t+1}$	

- Sample experiences from data-set and apply update

$$l = \left(r + \gamma \max_{a'} Q(s', a', \mathbf{w}^-) - Q(s, a, \mathbf{w}) \right)^2$$

- To deal with non-stationarity, target parameters \mathbf{w}^- are held fixed

Remember: Experience Replay

- Given **experience** consisting of $\langle \text{state}, \text{value} \rangle$, or $\langle \text{state}, \text{action/value} \rangle$ pairs

$$\mathcal{D} = \{ \langle s_1, v_1^\pi \rangle, \langle s_2, v_2^\pi \rangle, \dots, \langle s_T, v_T^\pi \rangle \}$$

- Repeat
 - Sample state, value from experience

$$\langle s, v^\pi \rangle \sim \mathcal{D}$$

- Apply stochastic gradient descent update

$$\Delta \mathbf{w} = \alpha (v^\pi - \hat{v}(s, \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(s, \mathbf{w})$$


DQNs: Experience Replay

- DQN uses experience replay and fixed Q-targets
- Store transition $(s_t, a_t, r_{t+1}, s_{t+1})$ in replay memory D
- Sample **random mini-batch** of transitions (s, a, r, s') from D
- Compute Q-learning targets w.r.t. old, fixed parameters w^-
- Optimize MSE between Q-network and Q-learning targets

$$\mathcal{L}_i(w_i) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}_i} \left[\underbrace{\left(r + \gamma \max_{a'} Q(s', a'; w_i^-) \right)}_{\text{Q-learning target}} - \underbrace{Q(s, a; w_i)}_{\text{Q-network}} \right]^2$$


- Use stochastic gradient descent

DQNs in Atari



DQNs in Atari


- ▶ End-to-end learning of values $Q(s,a)$ from pixels s
- ▶ Input state s is stack of raw pixels from last 4 frames
- ▶ Output is $Q(s,a)$ for 18 joystick/button positions
- ▶ Reward is change in score for that step



- ▶ Network architecture and hyperparameters fixed across all games

DQNs in Atari

- ▶ End-to-end learning of values $Q(s,a)$ from pixels s
- ▶ Input state s is stack of raw pixels from last 4 frames
- ▶ Output is $Q(s,a)$ for 18 joystick/button positions
- ▶ Reward is change in score for that step




DQN source code:
sites.google.com/a/deepmind.com/dqn/

- ▶ Network architecture and hyperparameters fixed across all games

Demo

DQN Results in Atari



Double Q-Learning

- Train 2 action-value functions, Q_1 and Q_2
- Do Q-learning on both, but
 - never on the same time steps (Q_1 and Q_2 are independent)
 - pick Q_1 or Q_2 at random to be updated on each step
- If updating Q_1 , use Q_2 for the value of the next state:

$$Q_1(S_t, A_t) \leftarrow Q_1(S_t, A_t) + \alpha \left(R_{t+1} + Q_2(S_{t+1}, \arg \max_a Q_1(S_{t+1}, a)) - Q_1(S_t, A_t) \right)$$

- Action selections are ε -greedy with respect to the sum of Q_1 and Q_2

Double Q-Learning

Initialize $Q_1(s, a)$ and $Q_2(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily

Initialize $Q_1(\text{terminal-state}, \cdot) = Q_2(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

 Initialize S

 Repeat (for each step of episode):

 Choose A from S using policy derived from Q_1 and Q_2 (e.g., ϵ -greedy in $Q_1 + Q_2$)

 Take action A , observe R, S'

 With 0.5 probability:

$$Q_1(S, A) \leftarrow Q_1(S, A) + \alpha \left(R + \gamma Q_2(S', \arg \max_a Q_1(S', a)) - Q_1(S, A) \right)$$

 else:

$$Q_2(S, A) \leftarrow Q_2(S, A) + \alpha \left(R + \gamma Q_1(S', \arg \max_a Q_2(S', a)) - Q_2(S, A) \right)$$

$S \leftarrow S'$;

 until S is terminal

Double DQN


- Current Q-network w is used to **select** actions
- Older Q-network w^- is used to **evaluate** actions

Action evaluation: w^-

$$l = \left(r + \gamma \overbrace{Q(s', \underbrace{\operatorname{argmax}_{a'} Q(s', a', w)}_{\text{Action selection: } w}), w^-}^{\text{Action evaluation: } w^-} - Q(s, a, w) \right)^2$$

Action selection: w

Double DQN



Prioritized Replay

- Weight experience according to **surprise**
- Store experience in priority queue according to DQN error

$$\left| r + \gamma \max_{a'} Q(s', a', \mathbf{w}^-) - Q(s, a, \mathbf{w}) \right|$$

- Stochastic Prioritization**

p_i is proportional to
DQN error

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha}$$

- α determines how much prioritization is used, with $\alpha = 0$ corresponding to the **uniform case**.

Dueling Networks


- Split Q-network into two channels
- **Action-independent** value function $V(s, v)$
- **Action-dependent** advantage function $A(s, a, w)$

$$Q(s, a) = V(s, v) + A(s, a, \mathbf{w})$$


- **Advantage function** is defined as:

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s).$$

Dueling Networks vs. DQN



DQN

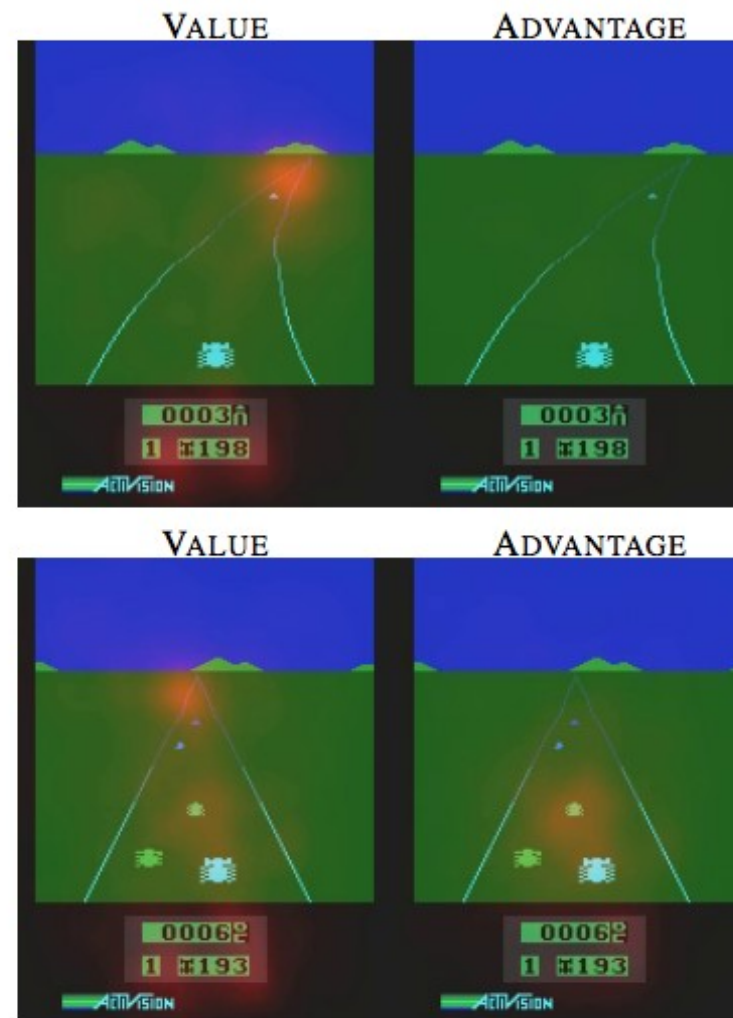


Dueling Networks


$$Q(s, a) = V(s, v) + A(s, a, \mathbf{w})$$

Dueling Networks

- The **value stream** learns to pay attention to the road
- The **advantage stream**: pay attention only when there are cars immediately in front, so as to avoid collisions



Dueling Networks



Multitask DQNs

- Can we train a single DQN to play multiple games at the same time



(Parisotto, Ba, Salakhutdinov, ICLR 2016)

Transfer Learning

- Can the network learn new games faster by leveraging knowledge about the previous games it learned.

